

YouTubeNxtWave learning p...W3Schools Online...CanvaStream | Google AL...LeetCodeMeta AIGoogle GeminiChatGPTPerplexityCourses | BrilliantNotebookLM

qlora_finetune.ipynbSave failedFile Edit View Insert Runtime Tools Help

CommandsCodeTextRun all

Connecting

Files

Connecting to a runtime to enable file browsing.

```
print(verbose, "\n", "end=")
_ = model.generate(
    input_ids = inputs,
    attention_mask = torch.ones_like(inputs), # FIXED: Removes the attention mask warning
    max_new_tokens = 256,
    use_cache = True,
    streamer = streamer # Streams text directly to the screen
)
print("\n" + "-" * 50)
```

Chat with your fine-tuned Macha AI (Type 'quit' to exit)

You: hi macha

Both 'max_new_tokens' (=256) and 'max_length' (=32768) seem to have been set. 'max_new_tokens' will take precedence. Please refer to the documentation for more information. (<https://h>)

Macha AI: Hey bro! What's up? How's your day going?<|im_end|>

You: hlo

Both 'max_new_tokens' (=256) and 'max_length' (=32768) seem to have been set. 'max_new_tokens' will take precedence. Please refer to the documentation for more information. (<https://h>)


Macha AI: Hello! How can I help you today?<|im_end|>

You: Bro explain what an API is

Both 'max_new_tokens' (=256) and 'max_length' (=32768) seem to have been set. 'max_new_tokens' will take precedence. Please refer to the documentation for more information. (<https://h>)

Macha AI: Macha, an API is short for Application Programming Interface, which allows different software applications to communicate with each other, mate.<|im_end|>

You:



portfolio_rag.ipynb

File Edit View Insert Runtime Tools Help

Commands Code Text Run all

Connected

3. Create FAISS index (Vector Database)

dimension = embeddings.shape[1]

index = faiss.IndexFlatL2(dimension)

index.add(embeddings)

print("Vector Database created successfully!")

Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

modules.json: 100% 349/349 [00:00<00:00, 31.8kB/s]

config_sentence_transformers.json: 100% 116/116 [00:00<00:00, 8.62kB/s]

README.md: 100% 10.5k/10.5k [00:00<00:00, 830kB/s]

sentence_bert_config.json: 100% 53.0/53.0 [00:00<00:00, 4.36kB/s]

config.json: 100% 612/612 [00:00<00:00, 44.1kB/s]

model.safetensors: reconstructing file: 100% 90.9MB / 90.9MB, 8.42MB/s

model.safetensors: downloading bytes: 85.0MB, 7.59MB/s

Loading weights: 100% 103/103 [00:00<00:00, 2684.50kB/s]

tokenizer_config.json: 100% 350/350 [00:00<00:00, 31.5kB/s]

vocab.txt: 100% 232k/232k [00:00<00:00, 8.50MB/s]

tokenizer.json: 100% 466k/466k [00:00<00:00, 13.7MB/s]

special_tokens_map.json: 100% 112/112 [00:00<00:00, 5.24kB/s]

config.json: 100% 190/190 [00:00<00:00, 15.4kB/s]

Vector Database created successfully!

5. The RAG Search Function

When you ask a question, this finds the most relevant information and adds it to the AI

Variables Terminal

15:06 Python 3