

USECASE 1 IHL

LLM

```
!pip install -q transformers accelerate

import torch

from transformers import pipeline

chatbot = pipeline(
    "text-generation",
    model="Qwen/Qwen2.5-0.5B-Instruct",
    device=0
)

print("🤖 Chatbot started!")

print("Type 'exit' to stop.\n")

while True:
    user_input = input("You: ")

    if user_input.lower() == "exit":
        print("Bot: Goodbye!")
        break

    messages = [
        {
            "role": "system",
            "content": "You are a helpful and friendly AI assistant."
        },
        {
            "role": "user",
            "content": user_input
        }
    ]
```

```
]
response = chatbot(
    messages,
    max_new_tokens=30,
    do_sample=True,
    temperature=0.7
)
answer = response[0]["generated_text"][-1]["content"]
print("Bot:", answer)
```

INPUT / OUTPUT:

```
You: WHAT COURSE DOES THE COLLEGE OFFER?
[transformers] Passing `generation_config` together with generation-related arguments=({'do_sample', 'ma
[transformers] Both `max_new_tokens` (=30) and `max_length` (=20) seem to have been set. `max_new_tokens`
[transformers] Ignoring clean_up_tokenization_spaces=True for BPE tokenizer Qwen2Tokenizer. The clean_up
Bot: The college offers various courses to suit different interests, career goals, and academic needs. H

1. General Education
You: EXIT
Bot: Goodbye!
```